

BASICS OF SQL

1. What is SQL and its significance in data analysis?

SQL (**Structured Query Language**) is a standard language for **storing, retrieving, and managing** data in databases. It helps analysts **extract insights** from large datasets efficiently.

Use Cases in Data Analysis:

- Filtering and grouping data
- Performing aggregations (SUM(), AVG(), etc.)
- Joining multiple tables
- Handling **big data** with optimized queries

2. Differentiate between SQL and MySQL.

Feature	SQL	MySQL
Definition	A query language	A database management system (DBMS)
Purpose	Used to manage relational databases	Implements SQL to manage data
Vendor	Standardized language	Owned by Oracle Corporation
Scalability	Used in various DBMS	Suitable for small to medium applications
Example Usage	SELECT * FROM employees;	MySQL databases store and execute SQL queries

3. Explain SQL joins: INNER JOIN, LEFT JOIN, RIGHT JOIN, and FULL JOIN.

Joins are used to **combine** data from two or more tables.

INNER JOIN (Returns only matching rows)

sql

```
SELECT employees.name, departments.department_name
FROM employees
INNER JOIN departments ON employees.department_id = departments.id;
```

LEFT JOIN (Returns all rows from the left table + matching ones from the right)

sql

```
SELECT employees.name, departments.department_name
FROM employees
LEFT JOIN departments ON employees.department_id = departments.id;
```

RIGHT JOIN (Returns all rows from the right table + matching ones from the left)

sql

```
SELECT employees.name, departments.department_name
FROM employees
RIGHT JOIN departments ON employees.department_id = departments.id;
```

FULL JOIN (Returns all rows when there is a match in either table)

sql

```
SELECT employees.name, departments.department_name
FROM employees
FULL JOIN departments ON employees.department_id = departments.id;
```

4. What are the primary components of a SQL query?

A typical SQL query includes:

1. **SELECT** – Specifies columns
2. **FROM** – Specifies table

3. **WHERE** – Filters records
4. **GROUP BY** – Groups results
5. **HAVING** – Filters grouped data
6. **ORDER BY** – Sorts results

Example Query:

sql

```
SELECT department, COUNT(*)  
FROM employees  
WHERE salary > 50000  
GROUP BY department  
HAVING COUNT(*) > 5  
ORDER BY department ASC;
```

5. Define the terms: table, row, and column in SQL.

- **Table:** A collection of related data in rows and columns.
- **Row (Record):** A single **entry** in a table.
- **Column (Field):** A specific attribute in a table.

Example:

ID	Name	Salary
1	Alice	50000
2	Bob	60000

Here,

- "employees" is the **table**
- (1, 'Alice', 50000) is a **row**
- "Name" is a **column**

6. How do you comment out lines in SQL?

- **Single-line comment:** -- This is a comment

- **Multi-line comment:** /* This is a multi-line comment */

7. What is the purpose of the **SELECT** statement?

The **SELECT** statement retrieves **specific data** from a table.

sql

```
SELECT name, salary FROM employees;
```

8. How do you retrieve all columns from a table?

Use * (wildcard):

sql

```
SELECT * FROM employees;
```

9. What is a **WHERE** clause used for?

Filters records based on conditions.

sql

```
SELECT * FROM employees WHERE salary > 50000;
```

10. Explain the difference between **WHERE** and **HAVING**.

Feature	WHERE	HAVING
Used for	Row filtering	Group filtering
Works with	SELECT, UPDATE, DELETE	GROUP BY
Example	WHERE salary > 50000	HAVING COUNT(*) > 5

11. How do you eliminate duplicate records?

Use DISTINCT:

sql

```
SELECT DISTINCT department FROM employees;
```

12. Difference between COUNT(*) and COUNT(column_name)

- COUNT(*) counts **all rows**, including NULLs.
- COUNT(column_name) counts **non-NULL values**.

13. GROUP BY vs ORDER BY

- GROUP BY: Groups rows by a column
- ORDER BY: Sorts the result

sql

```
SELECT department, COUNT(*) FROM employees GROUP BY department ORDER BY department;
```

14. How do you limit the number of records?

- **MySQL**: LIMIT 10
- **SQL Server**: TOP 10
- **Oracle**: FETCH FIRST 10 ROWS ONLY

15. Purpose of the LIKE operator?

Used for pattern matching.

sql

```
SELECT * FROM employees WHERE name LIKE 'A%'; -- Names starting with A
```

16. What is a wildcard character?

- % matches **zero or more** characters
- _ matches **exactly one** character

sql

```
SELECT * FROM employees WHERE name LIKE '_ohn'; -- Matches John, Mohn
```

17. How do you perform arithmetic operations?

sql

```
SELECT name, salary * 1.1 AS NewSalary FROM employees;
```

18. What is a subquery, and how is it used?

A subquery is a query inside another query.

sql

```
SELECT name FROM employees WHERE salary > (SELECT AVG(salary) FROM employees);
```

19. Purpose of the IN operator?

Checks if a value exists in a list.

sql

```
SELECT * FROM employees WHERE department_id IN (1, 2, 3);
```

20. Difference between UNION and UNION ALL

- UNION removes duplicates
- UNION ALL includes duplicates

sql

```
SELECT name FROM employees WHERE department_id = 1
```

```
UNION ALL
```

```
SELECT name FROM employees WHERE department_id = 2;
```

INTERMEDIATE SQL

1. What are aggregate functions in SQL? Provide examples.

Aggregate functions perform **calculations on multiple rows** and return a **single result**.

Common Aggregate Functions:

- SUM(): Adds values
- AVG(): Finds the average
- COUNT(): Counts rows
- MAX(): Finds the maximum
- MIN(): Finds the minimum

Example:

sql

```
SELECT department, COUNT(*) AS total_employees, AVG(salary) AS avg_salary
```

```
FROM employees
```

```
GROUP BY department;
```

2. How do you handle NULL values in SQL?

NULL means "no value."

Ways to handle NULL values:

- Use IS NULL or IS NOT NULL:

sql

```
SELECT * FROM employees WHERE salary IS NULL;
```

- Use COALESCE() to replace NULL with a default value:

sql

```
SELECT name, COALESCE(salary, 0) AS salary FROM employees;
```

3. INNER JOIN vs OUTER JOIN

JOIN Type	Description
INNER JOIN	Returns only matching rows
LEFT JOIN	Returns all left-table rows, and matching right-table rows
RIGHT JOIN	Returns all right-table rows, and matching left-table rows
FULL JOIN	Returns all records from both tables

4. What is a self-join in SQL?

A self-join joins a table to itself.

Example (Find employees with the same manager):

sql

```
SELECT e1.name AS Employee, e2.name AS Manager  
FROM employees e1
```


JOIN employees e2 ON e1.manager_id = e2.id;

5. How do you create a new table in SQL?

sql

```
CREATE TABLE employees (  
  id INT PRIMARY KEY,  
  name VARCHAR(100),  
  salary DECIMAL(10,2),  
  department_id INT  
);
```

6. What is a primary key, and why is it important?

A **primary key** uniquely identifies each row in a table.

- Ensures **data integrity**
- Prevents **duplicate records**

sql

```
CREATE TABLE students (  
  student_id INT PRIMARY KEY,  
  name VARCHAR(100)  
);
```

7. ACID properties in databases

ACID ensures **reliable transactions** in SQL:

- **Atomicity**: All or nothing
- **Consistency**: Data remains valid
- **Isolation**: Transactions are independent

- **Durability:** Data persists after a transaction

8. What is normalization, and why is it important?

Normalization **reduces redundancy** and **improves efficiency**.

Forms of Normalization:

- **1NF:** No duplicate columns
- **2NF:** No partial dependencies
- **3NF:** No transitive dependencies

9. How do you add new data to a table?

sql

```
INSERT INTO employees (id, name, salary, department_id)
VALUES (1, 'Alice', 60000, 2);
```

10. What is a stored procedure, and how is it created?

A stored procedure is a reusable SQL block.

sql

```
CREATE PROCEDURE GetEmployees()
AS
BEGIN
    SELECT * FROM employees;
END;
```

Run it using:

sql

EXEC GetEmployees;

11. View vs Table

Feature	Table	View
Data	Stores data	Stores query results
Storage	physically	dynamically
Updatable?	Yes	Sometimes
Use Case	Permanent storage	Read-only queries

Create a view:

sql

```
CREATE VIEW HighSalaryEmployees AS  
SELECT * FROM employees WHERE salary > 70000;
```

12. How do you modify an existing table structure?

- **Add a column:**

sql

```
ALTER TABLE employees ADD age INT;
```

- **Modify a column:**

sql

```
ALTER TABLE employees MODIFY COLUMN salary DECIMAL(12,2);
```

13. What is a transaction in SQL?

A **transaction** is a group of SQL statements executed together.

sql

```
BEGIN TRANSACTION;  
UPDATE employees SET salary = salary + 5000 WHERE id = 1;  
COMMIT;
```

14. Clustered vs Non-Clustered Index

Index Type	Description
Clustered Index	Data is physically sorted (only one per table)
Non-Clustered Index	Pointer-based indexing (multiple allowed)

sql

```
CREATE CLUSTERED INDEX idx_emp_salary ON employees(salary);  
CREATE NONCLUSTERED INDEX idx_emp_name ON employees(name);
```

15. How do you handle data concurrency in SQL?

- **Locks** prevent multiple transactions from modifying the same data.
- Use READ COMMITTED isolation level:

sql

```
SET TRANSACTION ISOLATION LEVEL READ COMMITTED;
```

16. Purpose of the EXPLAIN statement

EXPLAIN shows **query execution plans** to optimize performance.

sql

EXPLAIN SELECT * FROM employees WHERE salary > 50000;

17. What is a trigger, and when would you use one?

A trigger is an **automatic event** that runs before/after a change.

sql

```
CREATE TRIGGER update_salary
AFTER UPDATE ON employees
FOR EACH ROW
BEGIN
    INSERT INTO salary_log(employee_id, old_salary, new_salary)
    VALUES (OLD.id, OLD.salary, NEW.salary);
END;
```

18. UNION vs JOIN

Feature	UNION	JOIN
Combines	Rows from different queries	Columns from different tables
Removes		
Duplicates	Yes (UNION) / No (UNION ALL)	No
?		
Example	SELECT name FROM A UNION SELECT name FROM B;	SELECT A.name, B.salary FROM A JOIN B ON A.id = B.id;

19. How do you perform a case-insensitive search?

Use LOWER() or ILIKE (PostgreSQL).

sql

```
SELECT * FROM employees WHERE LOWER(name) = 'john';
```

20. Indexing in SQL and its advantages

Indexes **improve query performance**.

sql

```
CREATE INDEX idx_salary ON employees(salary);
```

ADVANCED SQL

1. What are Common Table Expressions (CTEs)?

CTEs **store temporary query results**.

sql

```
WITH HighSalary AS (  
    SELECT * FROM employees WHERE salary > 70000  
)  
SELECT * FROM HighSalary;
```

2. OLTP vs OLAP

Feature	OLTP	OLAP
Purpose	Transactions	Analytical Queries
Speed	Fast	Slow
Example	Bank transactions	Data warehouses

3. What is a window function?

Used for **ranking, cumulative sums, and moving averages**.

sql

```
SELECT name, salary,  
RANK() OVER (ORDER BY salary DESC) AS salary_rank  
FROM employees;
```

4. Purpose of COALESCE()

Replaces NULL values.

sql

```
SELECT name, COALESCE(salary, 0) AS salary FROM employees;
```

5. How to optimize SQL queries?

- Use **indexes**
- Avoid SELECT *
- Use **EXPLAIN** to analyze queries

6. Database Normalization (1NF, 2NF, 3NF)

- **1NF**: Atomic columns
- **2NF**: No partial dependencies
- **3NF**: No transitive dependencies

7. Correlated vs Non-Correlated Subquery

- **Correlated**: Depends on outer query
- **Non-Correlated**: Independent

sql

```
SELECT name FROM employees WHERE salary > (SELECT AVG(salary) FROM employees);
```

8. How do you perform data migration in SQL?

Data migration involves **transferring data** from one database to another. This can be done using **SQL commands**, ETL tools, or database migration services.

Steps for SQL-Based Data Migration:

1. **Extract Data** (from the source database):

sql

```
SELECT * INTO OUTFILE '/path/data.csv' FROM old_database.employees;
```

2. **Transform Data** (if needed, clean or modify data structure).
3. **Load Data** (into the target database):

sql

```
LOAD DATA INFILE '/path/data.csv' INTO TABLE new_database.employees;
```

4. **Verify Data** (compare counts and checksums between source and target).
5. **Optimize Performance** (add indexes, remove duplicates).

Database Denormalization and When to Use It

Denormalization **adds redundancy** to speed up queries.

When to Use Denormalization?

- **Read-heavy workloads** (e.g., analytics dashboards).
- **Fewer JOINS needed** (complex queries with many joins slow performance).
- **Precomputed aggregates** (like total sales, count of users).

Example: Denormalized Table

Instead of using **separate orders and customer tables**, store everything in **one table** for faster access:

sql

```
CREATE TABLE order_summary (  
  order_id INT PRIMARY KEY,  
  customer_name VARCHAR(255),  
  total_amount DECIMAL(10,2)  
);
```

9. Steps to Troubleshoot a Slow-Running SQL Query

1. Use EXPLAIN to Analyze Query Execution Plan

sql

```
EXPLAIN SELECT * FROM employees WHERE salary > 50000;
```

- a. Check for **full table scans** (bad)
- b. Look for **indexed columns**

2. Check for Missing Indexes

sql

```
CREATE INDEX idx_salary ON employees(salary);
```

3. Avoid SELECT * (Fetch Only Needed Columns)

sql

```
SELECT name, salary FROM employees WHERE salary > 50000;
```

4. Optimize WHERE Clause

- a. Use **indexed columns** in WHERE
 - b. Avoid **functions on columns**
 - c. Rewrite queries efficiently
- 5. Use Joins Efficiently**
- a. Prefer INNER JOIN over OUTER JOIN if possible
 - b. Ensure **JOINS use indexed columns**
- 6. Limit Results If Possible**

sql

```
SELECT * FROM orders LIMIT 100;
```

7. Analyze Locks and Transactions

sql

```
SHOW PROCESSLIST;
```

8. Check Query Execution Time

sql

```
SET profiling = 1;  
SELECT * FROM employees WHERE salary > 50000;  
SHOW PROFILES;
```

10. Steps to Troubleshoot a Slow-Running SQL Query

A slow SQL query can be caused by **poor indexing, inefficient joins, large data scans, or locking issues**. Here's a step-by-step approach to troubleshoot and optimize it:

1. Use EXPLAIN to Analyze Query Execution Plan

- The EXPLAIN statement helps identify how SQL executes a query.

- It shows **table scans, index usage, JOIN strategies, and cost estimation.**

Example:

sql

```
EXPLAIN SELECT * FROM employees WHERE salary > 50000;
```

What to check?

- **"Using Index"** (good) vs. **"Full Table Scan"** (bad)
- Look for **large row scans**

2. Identify Missing or Inefficient Indexes

Indexes **speed up search operations.**

- Use indexes on columns in WHERE, JOIN, ORDER BY.
- Avoid **over-indexing**, as it can slow down inserts/updates.

Example: Creating an Index on salary

sql

```
CREATE INDEX idx_salary ON employees(salary);
```

3. Optimize SELECT Queries (Avoid SELECT *)

Fetching **only necessary columns** speeds up queries.

Bad:

sql

```
SELECT * FROM employees WHERE department_id = 3;
```

Good:

sql

```
SELECT name, salary FROM employees WHERE department_id = 3;
```

4. Optimize WHERE Clause Usage

- Use **indexed columns** in the WHERE clause.
- **Avoid functions on indexed columns**, as they prevent index usage.

Bad (index won't work):

sql

```
SELECT * FROM employees WHERE YEAR(join_date) = 2020;
```

Good (index works):

sql

```
SELECT * FROM employees WHERE join_date BETWEEN '2020-01-01' AND '2020-12-31';
```

5. Optimize Joins (Use Proper Indexing)

Joins are often **the biggest performance bottleneck**.

Bad (No index on foreign keys):

sql

```
SELECT e.name, d.department_name  
FROM employees e  
JOIN departments d ON e.department_id = d.id;
```

Good (Indexing foreign keys):

sql

```
CREATE INDEX idx_department_id ON employees(department_id);
```

6. Use LIMIT to Reduce Result Set

If only a few records are needed, **limit the output** to avoid unnecessary data fetching.

sql

```
SELECT * FROM orders ORDER BY order_date DESC LIMIT 100;
```

7. Check and Reduce Locking Issues

Long-running transactions can block queries. Use:

sql

```
SHOW PROCESSLIST;
```

- Identify **locked queries**.
- Use **shorter transactions** or COMMIT frequently.

8. Analyze Execution Time with Profiling

Use SQL profiling to measure query performance.

sql

```
SET profiling = 1;
```

```
SELECT * FROM employees WHERE salary > 50000;
```

SHOW PROFILES;

9. Optimize Temporary Tables and Subqueries

Using **joins instead of subqueries** often improves performance.

Bad:

sql

```
SELECT * FROM employees WHERE department_id IN (SELECT id FROM departments
WHERE name = 'HR');
```

Good:

sql

```
SELECT e.* FROM employees e
JOIN departments d ON e.department_id = d.id
WHERE d.name = 'HR';
```

10. Partition Large Tables for Faster Queries

Partitioning **splits large tables into smaller ones** to speed up queries.

sql

```
CREATE TABLE orders (
  id INT,
  order_date DATE,
  amount DECIMAL(10,2),
  PRIMARY KEY (id, order_date)
) PARTITION BY RANGE (YEAR(order_date)) (
  PARTITION p1 VALUES LESS THAN (2022),
  PARTITION p2 VALUES LESS THAN (2023)
```

);

Final Checklist for Query Optimization

- ✓ Use EXPLAIN to analyze query execution.
- ✓ Ensure proper **indexing** on WHERE, JOIN, and ORDER BY columns.
- ✓ Avoid SELECT *; fetch only required columns.
- ✓ Optimize joins and subqueries.
- ✓ Check for **locking issues**.
- ✓ Use LIMIT for large datasets.
- ✓ Partition large tables if needed.